

MANAV RACHNA UNIVERSITY



SUPERVISED LEARNING



Submitted to:
- Dr. Roshni
Saxena

Submitted by: -
Harsh Gill
2K23CSUN01209
AIML - 3 A

INDEX

S. No.	Name of the Program	Date
1.	Write a python code to demonstrate commands for numpy and pandas.	
2.	Write a python program to calculate mean square and mean absolute error.	
3.	Write a python program to calculate gradient descent of a machine learning model.	
4.	Prepare a linear regression model for predicting the salary of user based on number of years of experience.	
5.	Prepare a linear regression model for prediction of resale car price.	
6.	Prepare a Lasso and Ridge regression model for prediction of house price and compare it with linear regression model.	
7.	Prepare a decision tree model for Iris Dataset using Gini Index.	
8.	Prepare a decision tree model for Iris Dataset using entropy.	
9.	Prepare a naïve bayes classification model for prediction of purchase power of a user.	
10.	Prepare a naïve bayes classification model for classification of email messages into spam or not spam.	
11.	Prepare a model for prediction of prostate cancer using KNN Classifier.	
12.	Prepare a model for prediction of survival from Titanic Ship using Random Forest and compare the accuracy with other classifiers also.	

Program 1

Write a python code to demonstrate commands for numpy and pandas.

```
# Demonstrate numpy commands
# Import necessary libraries
import numpy as np

# Creating arrays with zeros a = np.zeros(3)      #
1D array of zeros print("Array a:", a)
print("Type of array a:", type(a)) print("Type of
elements in array a:", type(a[0]))

b = np.zeros(3, dtype=int)      # 1D array of zeros with integer type
print("Array b:", b) print("Type of array b:", type(b))
print("Type of elements in array b:", type(b[0]))

# Reshape example z = np.zeros(3)
print("Original Array: ", z)
print("Shape of Array: ", z.shape)
z.shape = (3, 1)      # Reshape array to 5x1
print("Reshaped Array:\n", z) print("Shape
of Reshaped Array: ", z.shape)

# Creating an array using linspace z =
np.linspace(1, 2, 5) print("Array created
using linspace: ", z)

# Accessing array elements with positive and negative indexing
print("Element at index 0: ", z[0]) print("Element at index -
3: ", z[-3])      print("Array elements from index 0 to 2: ",
z[0:2])

# Identity matrix i =
np.identity(2, dtype=int)
print("Identity Matrix:\n", i)

# Creating a 2D matrix in two different ways
z = np.zeros((2, 2))      # 2D array of zeros
print("2-D Array (method 1):\n", z)

y = np.array([[1, 2], [3, 4]])      # Manually defined 2D array
print("2-D Array (method 2):\n", y)

# Accessing elements with index
print("Element at (0,1): ", y[0, 1])
print("Element at (0,0): ", y[0, 0])

# Slicing in 2D arrays
print("Second row: ", y[1, :])
print("First column: ", y[:, 0])

H = np.array([[1, 2, 3], [4, 5, 6], [7, 8, 9]])
print("2-D Array:\n", H) print("First row:",
H[0, :]) print("Third row:", H[2, :])
print("First column across rows: ", H[:, 0])

# Access elements at specified indices x = np.linspace(2,
4, 5) indices = np.array((0, 2, 3)) print("Array x:", x)
print("Elements at specified indices(0,2,3): ", x[indices])

# Boolean array d = np.array([0, 1, 2, 0, 0], dtype=bool)      # Every non-zero is
True, 0 is False print("Boolean Array d:", d)

# Sorting and basic array statistics a = np.array([17,
11, 15, 19, 24, 28, 26, 37, 35, 40])
a.sort() print("Original
Array:", a) print("Sorted
Array:", a) print("Sum:",
a.sum()) print("Min:",
a.min()) print("Max:",
a.max()) print("Argmin
(index of min):", a.argmin())
print("Argmax (index of
max):", a.argmax())
print("Cumulative Sum:",
a.cumsum())
print("Cumulative
Product:", a.cumprod())
print("Mean:", a.mean())
```

```

print("Median:",
np.median(a))
print("Variance:", a.var())
print("Standard
Deviation:", a.std())
print("Searchsorted (insert
position for 25):",
a.searchsorted(25))

# Array arithmetic operations
a = np.array([1, 2, 3, 4]) b =
np.array([5, 6, 7, 8]) print("a +
b:", a + b) print("a * b:", a * b)
print("a + 10:", a + 10) print("a *
10:", a * 10)

# Matrix operations
X = np.array([[1, 2, 3], [4, 5, 6], [5, 6, 7]]) Y
= np.array([[7, 8, 9], [4, 8, 9], [6, 3, 5]])
print("X:\n", X) print("Y:\n", Y) print("X +
Y:\n", X + Y) print("X + 10:\n", X + 10) print("X
* Y:\n", X @ Y) # Matrix multiplication
print("Transpose of X:\n", X.T)

# Comparison and modifying elements
Z = np.array([2, 3]) X =
np.array([2, 3])
print("X == Z:", X == Z)
X[0] = 5
print("X == Z after modifying X:", X == Z)

```



Show hidden output

```

# Import necessary libraries from pandas import DataFrame, Series # Import Series
and DataFrame for convenience import pandas as pd import numpy as np

```

```

# Creating a Series with default index ser_1 = Series([1, 1, 2, -3, -5, 8, 13])
print("Series with default index:\n", ser_1) print("Values in series: ",
ser_1.values) # Display only the values of the series

```

```

# Creating a Series with a custom index ser_2 = Series([1, 1, 2, -
3, -5], index=['a', 'b', 'c', 'd', 'e']) print("Series 2:\n",
ser_2)

```

```

# Accessing elements in a Series using index and labels
print("ser_2[1] == ser_2[b]", ser_2[1] == ser_2["b"])
print(ser_2[['c', 'a', 'b']]) # Filter Series for
values greater than 0 ser_2[ser_2 > 0]
# Apply an operation on Series elements
ser_2 * 2 np.exp(ser_2)

```

```

# Create a Series from a dictionary dict_1 = {'foo': 100, 'bar':
200, 'baz': 300} ser_3 = Series(dict_1) # Custom index on Series
index = ['foo', 'bar', 'baz', 'qux'] ser_4 = Series(dict_1,
index=index) # Missing values become NaN

```

```

# Print Series print("Series
3:\n", ser_3) print("Series
4:\n", ser_4) # Check for null
values in Series
print("Null values in ser_4:\n", pd.isnull(ser_4)) # Arithmetic
operations between Series print("Sum of series 3 and 4:\n", ser_3 +
ser_4) # Setting names for the Series and index ser_4.name =
'foobarbaz' ser_4.index.name = 'label' print("Series 4 after setting
names for series and index:\n", ser_4)

```

```

# Create another Series with custom index ser = Series([10, 15,
18, 12, 20, 9], index=[5, 8, 12, 0, 1, 7]) # Access elements by
label or position using loc and iloc print("Accessing elements
by label or position: ") print(ser.loc[0:1])
print(ser.iloc[0:1]) print(ser.iloc[0]) print(ser.loc[0])

```

```

# Create a DataFrame with dictionaries data_1 =
{'state': ['VA', 'VA', 'VA', 'MD', 'MD'],
'year': [2012, 2013, 2014, 2015, 2016],
'pop': [5.0, 5.1, 5.2, 4.0, 4.1]} df_1 =
DataFrame(data_1)

```

```

# Access a column of the DataFrame
df_1['state']

```

```

# Find and print the series of prime numbers from 1 to 300
primes = []
for i in range(1, 301):
    if i > 1:
        for j in range(2, i // 2 + 1):
            if i % j == 0:
                break
        else:
            primes.append(i)
primes_series = pd.Series(primes)
print("Series of Primes:\n", primes_series)

# Generate Fibonacci numbers up to 100
a, b = 0, 1
fibonacci_nums = []
while a < 100:
    fibonacci_nums.append(a)
    a, b = b, a + b
fibonacci_series = Series(fibonacci_nums)
print("Fibonacci Series:\n", fibonacci_series)

# Prompt user for a list of 20 numbers l = [int(x) for x
in input("Enter 20 numbers: ").split()]

# Initialize min, max, and sum variables
min_val = l[0]
max_val = l[0]
sum_val = 0

# Calculate sum, min, and max manually
for i in l:
    sum_val += i
    if i < min_val:
        min_val = i
    if i > max_val:
        max_val = i

print("Sum:", sum_val)
print("Min:", min_val)
print("Max:", max_val)

# Manually inputting values in a list one by one and finding the sum
l = []
sum_val = 0
for i in range(1, 21):
    num = int(input("Enter number: "))
    l.append(num)
    sum_val += num
print("Sum:", sum_val)

```



Show hidden output

```

▶ Array a: [0. 0. 0.]
Type of array a: <class 'numpy.ndarray'>
Type of elements in array a: <class 'numpy.float64'>
Array b: [0 0 0]
Type of array b: <class 'numpy.ndarray'>
Type of elements in array b: <class 'numpy.int64'>
Original Array: [0. 0. 0.]
Shape of Array: (3,)
Reshaped Array:
[[0.]
 [0.]
 [0.]]
Shape of Reshaped Array: (3, 1)
Array created using linspace: [1.  1.25 1.5  1.75 2. ]
Element at index 0: 1.0
Element at index -3: 1.5
Array elements from index 0 to 2: [1.  1.25]
Identity Matrix:
[[1 0]
 [0 1]]
2-D Array (method 1):
[[0. 0.]
 [0. 0.]]
2-D Array (method 2):
[[1 2]
 [3 4]]
Element at (0,1): 2
Element at (0,0): 1
Second row: [3 4]
First column: [1 3]
2-D Array:
[[1 2 3]
 [4 5 6]
 [7 8 9]]
First row: [1 2 3]
Third row: [7 8 9]
First column across rows: [1 4 7]
Array x: [2.  2.5 3.  3.5 4. ]

```

```

▶ Array x: [2.  2.5 3.  3.5 4. ]
Elements at specified indices(0,2,3): [2.  3.  3.5]
Boolean Array d: [False True True False False]
Original Array: [11 15 17 19 24 26 28 35 37 40]
Sorted Array: [11 15 17 19 24 26 28 35 37 40]
Sum: 252
Min: 11
Max: 40
Argmin (index of min): 0
Argmax (index of max): 9
Cumulative Sum: [ 11  26  43  62  86 112 140 175 212 252]
Cumulative Product: [ 11  165  2805  53295  1279080  33256080  931170240  32590958400  1205865460800  48234618432000]
Mean: 25.2
Median: 25.0
Variance: 87.55999999999999
Standard Deviation: 9.357350052231668
Searchsorted (insert position for 25): 5
a + b: [ 6  8 10 12]
a * b: [ 5 12 21 32]
a + 10: [11 12 13 14]
a * 10: [10 20 30 40]
X:
[[1 2 3]
 [4 5 6]
 [5 6 7]]
Y:
[[7 8 9]
 [4 8 9]
 [6 3 5]]
X + Y:
[[ 8 10 12]
 [ 8 13 15]
 [11  9 12]]
X + 10:
[[11 12 13]
 [14 15 16]
 [15 16 17]]
X * Y:
[[ 33  33  42]
 [ 84  90 111]
 [101 109 134]]
Transpose of X:
[[1 4 5]
 [2 5 6]
 [3 6 7]]
[2 3]
X == Z: [ True  True]
X == Z after modifying X: [ True False]

```

```
print("ser_2[1] == ser_2[b]", ser_2[1] == ser_2["b"])
```

Series with default index:

```
0    1
1    1
2    2
3   -3
4   -5
5    8
6   13
```

dtype: int64

Values in series: [1 1 2 -3 -5 8 13]

Series 2:

```
a    1
b    1
c    2
d   -3
e   -5
```

dtype: int64

ser_2[1] == ser_2[b] True

```
c    2
```

```
a    1
```

```
b    1
```

dtype: int64

Series 3:

```
foo    100
bar    200
baz    300
```

dtype: int64

Series 4:

```
foo    100.0
bar    200.0
baz    300.0
```

```
qux     NaN
```

dtype: float64

Null values in ser_4:

```
foo    False
```

```
bar    False
```

```
baz    False
```

```
qux     True
```

dtype: bool

Sum of series 3 and 4:

```
bar    400.0
```

```
baz    600.0
```

```
foo    200.0
```

```
qux     NaN
```

dtype: float64

Series 4 after setting names for series and index:

```
label
```

```
foo    100.0
```

```
bar    200.0
```

```
baz    300.0
```

```
qux     NaN
```

Name: foobarbaz, dtype: float64

Accessing elements by label or position:

```
0    12
```

```
1    20
```

dtype: int64

```
5    10
```

dtype: int64

```
10
```

```
12
```

Series of Primes:

```
0    2
```

```
1    3
```

```
2    5
```

```
3    7
```

```
4   11
```

```
...
```

```
57  271
```

```
58  277
```

```
59  281
```

```
60  283
```

```
61  293
```

```
Length: 62, dtype: int64
Fibonacci Series:
0      0
1      1
2      1
3      2
4      3
5      5
6      8
7     13
8     21
9     34
10    55
11    89
dtype: int64
Enter 20 numbers: 1 2 3 4 5 6 7 8 0 9 11 23 44 2 21 34 5 12 23 21
Sum: 241
Min: 0
Max: 44
Enter number: 1
Enter number: 2
Enter number: 3
Enter number: 4
Enter number: 5
Enter number: 6
Enter number: 3
Enter number: 62
Enter number: 47
Enter number: 34
Enter number: 67
Enter number: 433
Enter number: 33
Enter number: 25
Enter number: 24
Enter number: 54
Enter number: 53
Enter number: 2
```


Program 2


Write a python program to calculate mean absolute error and mean square error.

```
#function to calculate the predicted values
def predicted_output(x,w,b):
    y_hat=[]
    for i in range(len(x)):
        y_hat.append(w*x[i]+b)
    return y_hat

#function to calculate mean absolute error
def MAE(y, y_hat):
    sum=0
    for i in range(len(y)):
        sum+=abs(y_hat[i]-y[i])
    return sum/len(y)

#function to calculate mean square error
def MSE(y, y_hat):
    sq_sum=0
    for i in range(len(y)):
        sq_sum+=(y_hat[i]-y[i])**2
    return sq_sum/len(y)

#taking inputs x=[eval(x) for x in input("Enter the values of x(input) separated by ',' :
").split(",")] y=[eval(x) for x in input("Enter the values of y(output) separated by ',' :
").split(",")] w=eval(input("Enter the value of w: ")) b=eval(input("Enter the value of
b: "))
#calling functions
y_hat=predicted_output(x, w, b)
MAE_value=MAE(y, y_hat)
MSE_value=MSE(y, y_hat) #printing the
values print("Predicted Output: ",y_hat)
print("Mean Absolute Error: ",MAE_value)
print("Mean Square Error: ",MSE_value)
```

```
 Enter the values of x(input) separated by ',': 3, 6, 9, 12, 15, 18, 20
Enter the values of y(output) separated by ',': 15, 28, 63, 90, 120, 152, 190
Enter the value of w: 2.5
Enter the value of b: 0

Predicted Output: [7.5, 15.0, 22.5, 30.0, 37.5, 45.0, 50.0]
Mean Absolute Error: 64.35714285714286
Mean Square Error: 6188.678571428572
```



Write a python program to calculate gradient descent of a machine learning model.

Program 3

```
# Import necessary libraries
import numpy as np
import matplotlib.pyplot as plt

# Function to perform gradient descent
def gradient_descent(func, x, learning_rate, num_iterations):
    x_values = []
    for i in range(num_iterations):
        gradient = func(x)
        x_values.append(x)
        x -= (learning_rate * gradient)
    return x, x_values

# Define the original function
def function(x):
    return 4*x**2 + 3*x + 1
# Define the derivative of the function
def derivative_f(x):
    return 8*x + 3

# Plotting the gradient descent steps on the function curve
def plot_gradient_descent(func, x, learning_rate, num_iterations, x_values):
    x_range = np.linspace(-10, 10, 400)
    y_range = func(x_range)
    plt.plot(x_range, y_range, label="f(x) = 4x^2 + 3x + 1")
    plt.scatter(x_values, [func(x) for x in x_values], color='red', label="Gradient Descent Steps")
    plt.xlabel("x")
    plt.ylabel("f(x)")
    plt.legend()
    plt.grid(True)
    plt.title("Gradient Descent Path on f(x)")
    plt.show()

# Set parameters for gradient descent
initial_x = 10
learning_rate = 0.01
num_iterations = 25

# Perform gradient descent
min_x, x_values = gradient_descent(derivative_f, initial_x, learning_rate, num_iterations)

# Print results
print("For function: 4x^2 + 3x + 1: ")
print("Minimum value of x:", min_x)

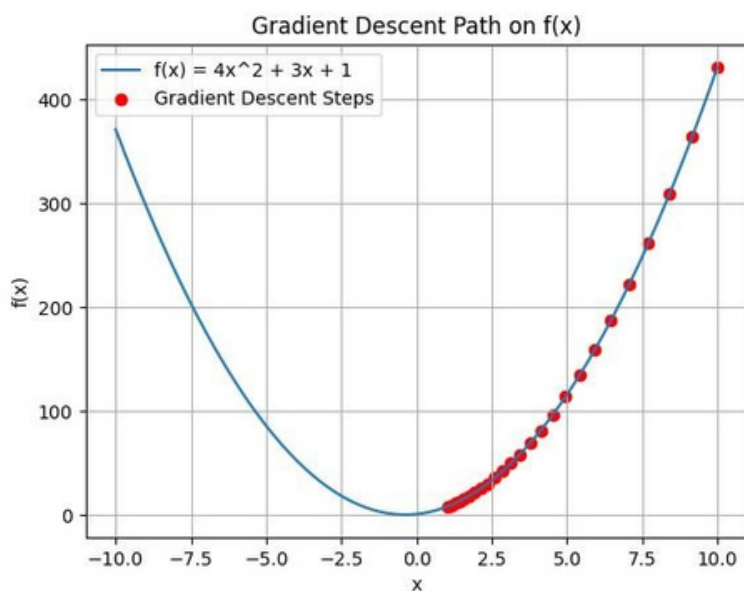
# Call the plot function to visualize gradient descent
plot_gradient_descent(function, x, learning_rate, num_iterations, x_values)

# Prepare a linear regression model for predicting the salary of user based on number of years of experience.
```

```
# importing necessary libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# loading the dataset df =
pd.read_csv('Salary_Data.csv')
```

↗ For function: $4x^2 + 3x + 1$:
Minimum value of x: 0.9152794755738098



🔗 Program 4

```
# defining the feature variable 'x' by dropping Salary and target variable 'y' as the Salary column
x = df.drop('Salary', axis=1) y = df['Salary']

# split the dataset into training and testing sets from sklearn.model_selection import
train_test_split x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3,
random_state=1)

# initialize and train the Linear Regression model on the training data
from sklearn.linear_model import LinearRegression model =
LinearRegression() model.fit(x_train, y_train)

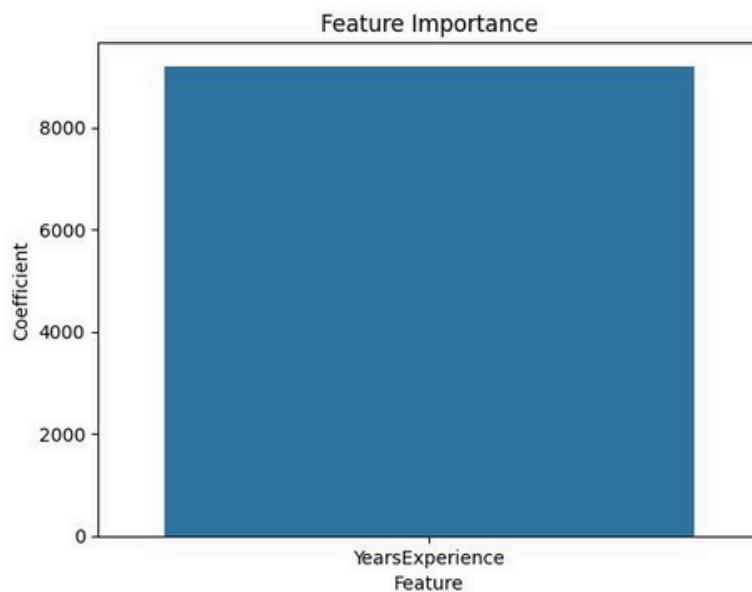

# Predict the target variable for the test set
y_test_predict = model.predict(x_test)

# Display the model's coefficient and intercept print("Model
coefficient(s):", model.coef_) print("Model intercept:",
model.intercept_) print("Model R^2 score on test set:",
model.score(x_test, y_test))

# scatter plot to visualize the relationship between predicted and actual values in the test set
plt.figure(figsize=(6, 3)) plt.scatter(y_test, y_test_predict, color='blue', label="Predicted vs
Actual") plt.xlabel("Actual Salary") plt.ylabel("Predicted Salary") plt.title("Actual vs
Predicted Salary (Test Set)") plt.legend() plt.show()

# bar plot to display the importance of each feature based on model coefficients
imp=pd.DataFrame(list(zip(x_test.columns,np.abs(model.coef_))),columns=['Feature','Coefficient'])
sns.barplot(x='Feature', y='Coefficient', data=imp) plt.title("Feature Importance") plt.show()
```

Modelcoefficient(s):[9202.23359825]
Modelintercept:26049.577715443353
ModelR^2scoreontestset:0.9248580247217075



Program 5

Prepare a linear regression model for prediction of resale car price.

```
# import necessary libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# load the dataset
df = pd.read_csv('cars24-car-price-cleaned.csv')

# replace 'make' and 'model' columns with the mean selling price for each group
df['make'] = df.groupby('make')['selling_price'].transform('mean')
df['model'] = df.groupby('model')['selling_price'].transform('mean')

# normalize the dataset using MinMaxScaler to scale features between 0 and 1
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
df_normalized = pd.DataFrame(scaler.fit_transform(df), columns=df.columns)

# define target variable 'y' as the selling price and features 'x' by dropping the selling price
y = df_normalized['selling_price']
x = df_normalized.drop('selling_price', axis=1)

# split the dataset into training and testing sets from sklearn.model_selection
import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=1)

# initialize and train the Linear Regression model on the training data
from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(x_train, y_train)
```



```
# predict the target variable for the test set
y_test_predict = model.predict(x_test)
```

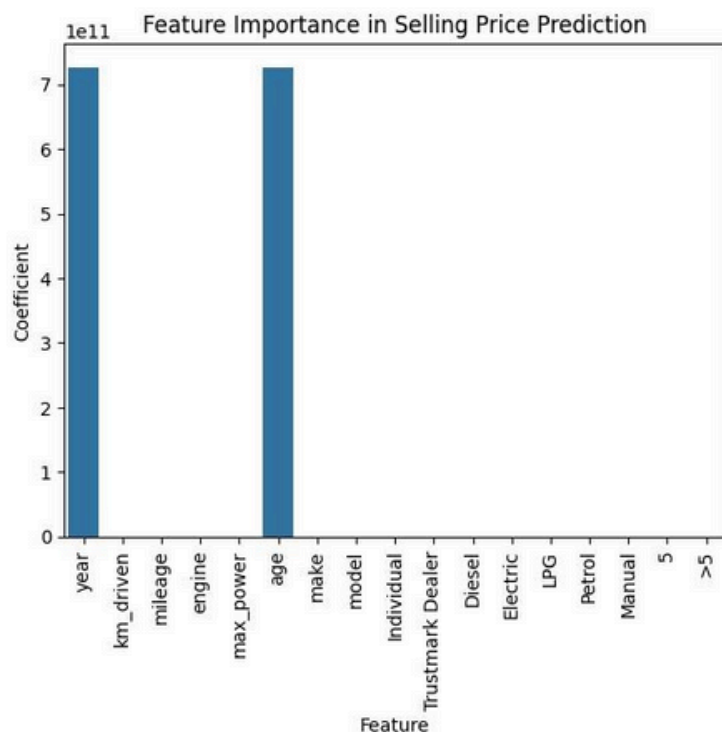

```
# Display model's coefficient, intercept, and R^2 score on test set
print("Model coefficients:", model.coef_) print("Model intercept:",
model.intercept_) print("Model R^2 score on test set:",
model.score(x_test, y_test))
```

```
# Scatter plot to visualize the relationship between predicted and actual values in the test set
#plt.figure(figsize=(8, 6)) plt.scatter(y_test, y_test_predict, label="Predicted vs Actual")
plt.xlabel("Actual Selling Price (Normalized)") plt.ylabel("Predicted Selling Price
(Normalized)") plt.title("Actual vs Predicted Selling Price (Test Set)") plt.legend() plt.show()
```

```
# Bar plot to display the importance of each feature based on model coefficients imp =
pd.DataFrame(list(zip(x_test.columns, np.abs(model.coef_))), columns=['Feature', 'Coefficient'])
#plt.figure(figsize=(8, 6)) sns.barplot(x='Feature',
y='Coefficient', data=imp) plt.xticks(rotation=90)
plt.title("Feature Importance in Selling Price Prediction")
plt.show()
```



```
Model coefficients: [ 7.26831852e+11 -2.50610352e-01 -2.32537818e-01      7.38776447e-02
 4.70141495e-02  7.26831852e+11  6.62815814e-02  8.59178586e-01 -
 7.22882618e-03 -7.02099753e-03  7.03528760e-03  1.32983308e-01
 1.49877118e-02 -6.86552095e-03 -3.59124005e-03 -1.61993065e-02 -
 2.35818239e-02]
Model intercept: -726831852169.8219
Model R^2 score on test set: 0.9459835819294395
```



Program 6

Prepare a Lasso and Ridge regression model for prediction of house price and compare it with linear regression model.

```
# Import necessary libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression, Lasso, Ridge
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import MinMaxScaler

# Load the housing dataset
df = pd.read_csv('Housing.csv')

# Convert categorical variables into numerical features that can be used by the model (target variable encoding)
df['mainroad'] = df.groupby('mainroad')['price'].transform('mean')
df['guestroom'] = df.groupby('guestroom')['price'].transform('mean')
df['basement'] = df.groupby('basement')['price'].transform('mean')
df['hotwaterheating'] = df.groupby('hotwaterheating')['price'].transform('mean')
df['airconditioning'] = df.groupby('airconditioning')['price'].transform('mean')
df['prefarea'] = df.groupby('prefarea')['price'].transform('mean')
df['furnishingstatus'] = df.groupby('furnishingstatus')['price'].transform('mean')
```



```
# Normalize the dataset to bring all features to the same scale scaler =
MinMaxScaler() df_normalized = pd.DataFrame(scaler.fit_transform(df),
columns=df.columns)

# Define the target variable 'y' as 'median_house_value' and features 'x' by dropping the target column
y = df_normalized['price'] x = df_normalized.drop('price', axis=1)

# Split the dataset into training and testing sets x_train, x_test, y_train, y_test =
train_test_split(x, y, test_size=0.3, random_state=1)

# Initialize models: Linear Regression, Lasso Regression, and Ridge Regression
model = LinearRegression() lasso_model = Lasso(alpha=0.1) ridge_model =
Ridge(alpha=0.1)

# Fit each model to the training data
model.fit(x_train, y_train)
lasso_model.fit(x_train, y_train)
ridge_model.fit(x_train, y_train)

# Display model coefficients, intercepts and R^2 scores
print("Linear Regression Coefficients:", model.coef_)
print("Lasso Regression Coefficients:", lasso_model.coef_)
print("Ridge Regression Coefficients:", ridge_model.coef_)

print("Linear Regression Intercept:", model.intercept_)
print("Lasso Regression Intercept:", lasso_model.intercept_)
print("Ridge Regression Intercept:", ridge_model.intercept_)

print("Linear Regression R^2 Score (Train):", model.score(x_train, y_train))
print("Lasso Regression R^2 Score (Train):", lasso_model.score(x_train, y_train))
print("Ridge Regression R^2 Score (Train):", ridge_model.score(x_train, y_train))


# Predict the target values on the test set using each model
y_pred = model.predict(x_test) y_pred_lasso =
lasso_model.predict(x_test) y_pred_ridge =
ridge_model.predict(x_test)

# Calculate Mean Squared Error (MSE) for each model on the test set
mse = mean_squared_error(y_test, y_pred) mse_lasso =
mean_squared_error(y_test, y_pred_lasso) mse_ridge =
mean_squared_error(y_test, y_pred_ridge)

# Display the MSE results to compare model performance, with lower MSE indicating better fit
print('MSE without regularization (Linear Regression):', mse) print('MSE with Lasso
regularization:', mse_lasso) print('MSE with Ridge regularization:', mse_ridge)

# Visualize the comparison of actual vs predicted values for each model
plt.figure(figsize=(10, 6)) plt.scatter(y_test, y_pred, color='blue', label="Linear
Regression Predictions") plt.scatter(y_test, y_pred_lasso, color='green', label="Lasso
Regression Predictions") plt.scatter(y_test, y_pred_ridge, color='red', label="Ridge
Regression Predictions", marker='*') plt.xlabel("Actual Price") plt.ylabel("Predicted
Price") plt.title("Comparison of Predictions by Different Regression Models")
plt.legend() plt.show()
```

 Linear Regression Coefficients: [0.31039697 0.01959006 0.26477477 0.13658528 0.04098972 0.02376751
0.04792801 0.07098812 0.05282266 0.07096655 0.04358941 0.03623753]
Lasso Regression Coefficients: [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
Ridge Regression Coefficients: [0.30639084 0.02106921 0.26241647 0.13615958 0.04133038 0.02401481
0.04774817 0.07051319 0.0530351 0.0713936 0.04377635 0.03640865]
Linear Regression Intercept: -0.0050427725675667445
Lasso Regression Intercept: 0.26192224608287595
Ridge Regression Intercept: -0.0048457449783638196
Linear Regression R^2 Score (Train): 0.6806547764599723
Lasso Regression R^2 Score (Train): 0.0
Ridge Regression R^2 Score (Train): 0.6806349211986238
MSE without regularization (Linear Regression): 0.010274158458096141
MSE with Lasso regularization: 0.03051838551799671
MSE with Ridge regularization: 0.010266744866035897



✓ Program 7

Prepare a decision tree model for Iris Dataset using Gini Index.

```
# Import necessary libraries
from sklearn import datasets
from sklearn.tree import DecisionTreeClassifier,
plot_tree from sklearn.metrics import accuracy_score
import matplotlib.pyplot as plt import pandas as pd

# Load the Iris dataset df
= pd.read_csv("Iris.csv")

# Define feature matrix 'x' by dropping 'Species' and 'Id' columns and target variable 'y' as
'Species' x = df.drop(['Species', 'Id'], axis=1) y = df['Species']

# Initialize DecisionTreeClassifier with Gini impurity criterion
model = DecisionTreeClassifier(criterion='gini')

# Dictionary to store Gini impurity for each feature
gini_impurities = {}

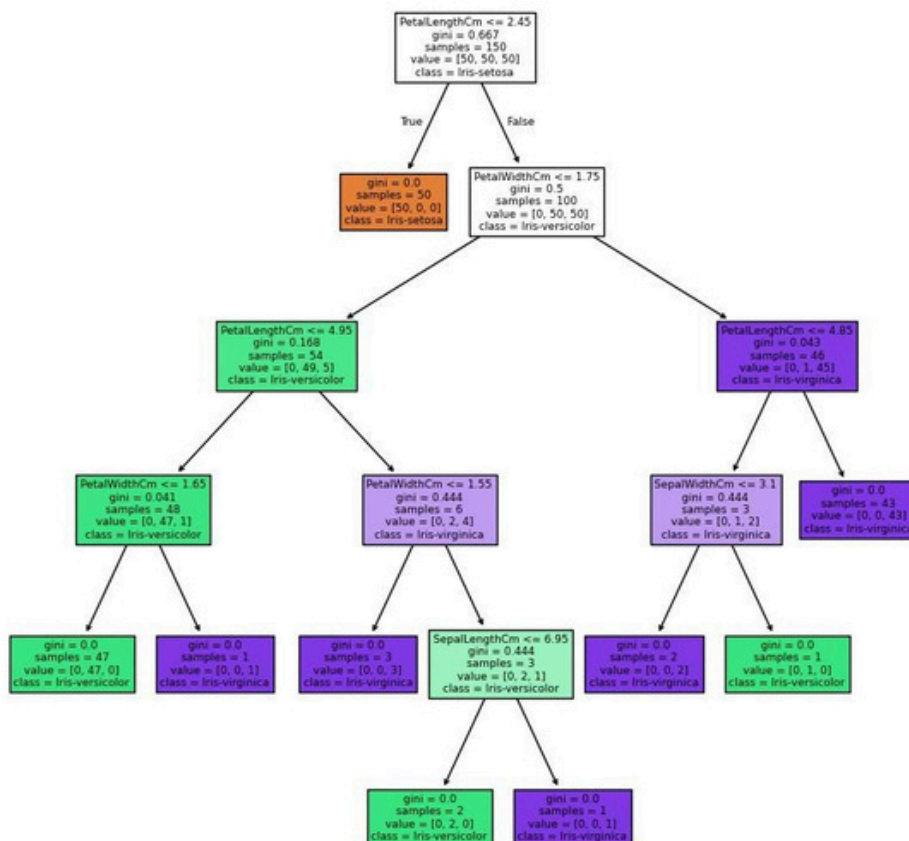
#loop through each feature
for i in range(x.shape[1]):
#fit classifier with only the current feature
model.fit(x.iloc[:, i].values.reshape(-1, 1), y)
prob=model.predict_proba(x.iloc[:, i].values.reshape(-1,1))
gini_impurities[i] = 1 - (prob[:, 0]**2 + prob[:, 1]**2 + prob[:, 2]**2).sum()

# Find the feature with the lowest Gini impurity (best
feature) best_feature = min(gini_impurities,
key=gini_impurities.get) print(f"Best feature:
{x.columns[best_feature]}") model.fit(x, y)

#plot original tree plt.figure(figsize=(10, 10)) plot_tree(model, filled=True,
feature_names=x.columns, class_names=model.classes_) plt.title("Original
Decision Tree")
plt.show()
```

➡ Bestfeature:PetalLengthCm

Original Decision Tree



✓ Program 8

Prepare a decision tree model for Iris Dataset using entropy.

```
# Import necessary libraries
import numpy as np
import pandas as pd
from sklearn.metrics import confusion_matrix, accuracy_score
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier, plot_tree
import matplotlib.pyplot as plt

# Load the Iris dataset
df = pd.read_csv("Iris.csv")

# Define feature matrix 'x' by dropping 'Species' and 'Id' columns and target variable 'y' as 'Species'
x = df.drop(["Species", "Id"], axis=1)
y = df["Species"]

# Splitting the dataset into train and test
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=100)

# Build decision tree
model = DecisionTreeClassifier(criterion='entropy', max_depth=4)

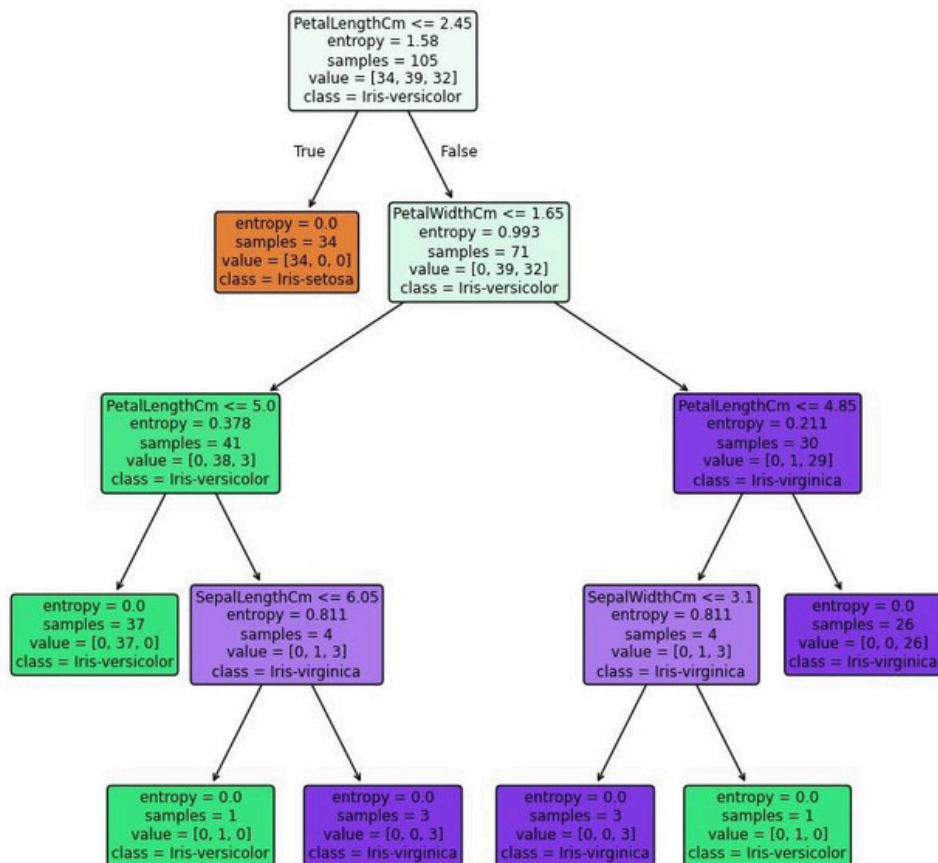
# Fit the tree to iris dataset
model.fit(x_train, y_train)

# Find the accuracy of the model
y_pred = model.predict(x_test)
print("Accuracy: ", accuracy_score(y_test, y_pred)*100)

# Function to plot the decision tree
def plot_decision_tree(model, feature_names, class_names):
    plt.figure(figsize=(10, 10))
    plot_tree(model, filled=True, feature_names=feature_names, class_names=class_names, rounded=True)
    plt.show()

plot_decision_tree(model, ["SepalLengthCm", "SepalWidthCm", "PetalLengthCm", "PetalWidthCm"], ["Iris-setosa", "Iris-versicolor", "Iris-virginica"])
```

Accuracy: 95.55555555555556



Program 9

Prepare a naïve bayes classification model for prediction of purchase power of a user.

```
# Import libraries import pandas as pd import numpy as np import matplotlib.pyplot as plt from matplotlib.colors import ListedColormap import seaborn as sns from sklearn.preprocessing import LabelEncoder, StandardScaler from sklearn.model_selection import train_test_split from sklearn.naive_bayes import GaussianNB from sklearn import metrics from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, precision_recall_curve, f1_score
```

```
# Load User_Data dataset df = pd.read_csv('User_Data.csv')
```

```
# Drop User ID column as it does not contribute towards prediction purpose df.drop(['User ID'], axis=1, inplace=True)
```

```
# Label Encoding le=LabelEncoder() df['Gender']=le.fit_transform(df['Gender'])
```

```
# Split data into dependent/independent variables x = df.iloc[:, :-1].values y = df.iloc[:, -1].values
```

```
# Split the dataset into training and testing sets x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.25, random_state=True)
```

```
# Scale dataset sc = StandardScaler() x_train = sc.fit_transform(x_train) x_test = sc.transform(x_test)
```

```
# Create naive-bayes classifier model classifier=GaussianNB() classifier.fit(x_train, y_train)
```

```
# Predict the values y_pred=classifier.predict(x_test) # Print accuracy of classifier print("Accuracy of classifier: ", accuracy_score(y_test, y_pred))
```

```
# Print the classification report print(f'Classification report:\n{classification_report(y_test, y_pred)}')
```


```
# Print the confusion matrix cf_matrix=confusion_matrix(y_test, y_pred) sns.heatmap(cf_matrix, annot=True, fmt='d', cmap='Blues', cbar=False)
```

Accuracy of classifier: 0.87

Classification report:

	recall	f1-score	support	precision
0	0.89	0.88	0.89	42
1	0.84	0.86	0.85	58
accuracy				0.87
macro avg	0.87	0.87	0.87	100
weighted avg	0.87	0.87	0.87	100

<Axes: >



Program 10

Prepare a naïve bayes classification model for classification of email messages into spam or not spam.

```
# Import libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import MultinomialNB, GaussianNB
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import accuracy_score, f1_score
import matplotlib.pyplot as plt
from wordcloud import WordCloud

# Load the dataset into a DataFrame with 'latin-1' encoding to avoid encoding issues
df = pd.read_csv('spam.csv', encoding='latin-1')

# Select only the relevant columns ('v1' as labels and 'v2' as messages) and rename them
df = df[['v1', 'v2']]
df = df.rename(columns={'v1': 'label', 'v2': 'text'})

# Define feature matrix 'x' as 'text' and target variable 'y' as 'label'
x = df['text']
y = df['label']

# Split the dataset into training and testing sets x_train, x_test, y_train, y_test =
train_test_split(x, y, test_size=0.2, random_state=42)

# Find and plot the distribution of spam and ham messages
distribution = y.value_counts()
print("Distribution of spam and ham messages:\n", distribution)
distribution.plot(kind='pie', autopct='%1.1f%%')
plt.title("Distribution of Spam and Ham Messages")
plt.show()

# Generate a Wordcloud for the Spam emails
spam_text = ' '.join(df[df['label'] == 'spam']['text'])
spam_wordcloud = WordCloud(width=800, height=400, max_words=100, background_color='white', random_state=42).generate(spam_text)

# Generate a Wordcloud for the Ham emails
ham_text = ' '.join(df[df['label'] == 'ham']['text'])
ham_wordcloud = WordCloud(width=800, height=400, max_words=100, background_color='white', random_state=42).generate(ham_text)

# Plot the word clouds for spam messages
plt.figure(figsize=(10, 4))
plt.subplot(1, 2, 1)
plt.imshow(spam_wordcloud)
plt.title('Word Cloud for Spam Messages')
plt.axis('off')

# Plot the wordcloud for ham messages
plt.subplot(1, 2, 2)
plt.imshow(ham_wordcloud)
plt.title('Word Cloud for Ham Messages')
plt.axis('off')

# Show both plots side by side
plt.tight_layout()
plt.show()

# Vectorize the text data to convert it into numerical features
vectorizer = CountVectorizer()
x_train = vectorizer.fit_transform(x_train)
x_test = vectorizer.transform(x_test)

# Train a Multinomial Naive Bayes classifier on the vectorized data
model_multinomial = MultinomialNB(alpha=0.8, fit_prior=True, force_alpha=True)
model_multinomial.fit(x_train, y_train)

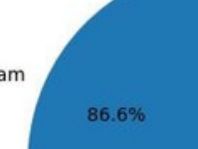
# Train a Gaussian Naive Bayes classifier on the vectorized data
model_gaussian = GaussianNB()
model_gaussian.fit(x_train.toarray(), y_train)

# Calculate and print the accuracy of both models on the test data
y_pred_multinomial = model_multinomial.predict(x_test)
accuracy_multinomial = accuracy_score(y_test, y_pred_multinomial)
print("Accuracy for Multinomial Naive Bayes Model: ", accuracy_multinomial)

y_pred_gaussian = model_gaussian.predict(x_test.toarray())
accuracy_gaussian = accuracy_score(y_test, y_pred_gaussian)
print("Accuracy for Gaussian Naive Bayes Model: ", accuracy_gaussian)

# Plot a comparison of the accuracy scores for the two classification methods
methods = ["Multinomial Naive Bayes", "Gaussian Naive Bayes"]
scores = [accuracy_multinomial, accuracy_gaussian]
plt.bar(methods, scores)
plt.xlabel("Classification Methods")
plt.ylabel("Accuracy")
plt.title("Comparison of Classification Methods")
plt.show()
```


```
➡ Distribution of spam and ham messages:  
label ham      4825 spam      747 Name:  
count, dtype: int64
```



Category	Count	Percentage
ham	86.6%	86.6%
spam	13.4%	13.4%

[illegible][illegible]

Accuracy for Multinomial Naive Bayes Model: 0.9838565022421525
Accuracy for Gaussian Naive Bayes Model: 0.9004484304932735



Classification Methods	Accuracy
Multinomial Naive Bayes	0.98
Gaussian Naive Bayes	0.90

Program 11

Prepare a model for prediction of prostate cancer using KNN Classifier.

```
# Import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
```

```

# Load the dataset df =
pd.read_csv('prostate.csv')

# Define feature matrix 'x' and target vector 'y'
x=df.drop('Target', axis = 1) y=df['Target']

# Feature scaling using StandardScaler scaler=StandardScaler()
df1=pd.DataFrame(scaler.fit_transform(x),columns=x.columns[::1])

# Split data into training and testing sets
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=1)

# Initialize K-Nearest Neighbors classifier with 1 neighbor
knn_model = KNeighborsClassifier(n_neighbors=1) knn_model.fit(x_train,y_train)

# Make predictions on the test set
y_pred = knn_model.predict(x_test)

# Display the confusion matrix to evaluate model performance
print("Confusion Matrix:\n", confusion_matrix(y_test,y_pred))

# Display classification report with precision, recall, F1-score, and accuracy
print("Classification Report:\n", classification_report(y_test,y_pred))

# Elbow method for determining the optimal number of neighbors 'K'
error_rate = []
for i in range(1,40):
    knn = KNeighborsClassifier(n_neighbors=i)
    knn.fit(x_train,y_train)
    new_y_pred = knn.predict(x_test)
    error_rate.append(np.mean(new_y_pred != y_test))

# Plot the error rate for different values of K
plt.figure(figsize=(12,5))
plt.plot(error_rate,color='blue',
linestyle='dashed', marker='o',
markerfacecolor='red', markersize=10)
plt.title('Error Rate vs. K Value')
plt.xlabel('K')
plt.ylabel('Error Rate')
plt.show()

```


➡ Confusion Matrix:

4][18
2][6

Classification Report:
recall f1-score support

precision

	0	0.75	0.82	0.78	22
1	0.33	0.25	0.29	8	
accuracy				0.67	30
macro avg	0.54	0.53	0.53		30
weighted avg	0.64	0.67	0.65		30





Program 12

Prepare a model for prediction of survival from Titanic Ship using Random Forest and compare the accuracy with other classifiers also.

```

# Import necessary libraries import pandas as pd from sklearn.model_selection
import train_test_split from sklearn.ensemble import RandomForestClassifier from
sklearn.metrics import accuracy_score, classification_report, confusion_matrix from
sklearn.preprocessing import LabelEncoder from sklearn.neighbors import
KNeighborsClassifier from sklearn.naive_bayes import GaussianNB from sklearn.tree
import DecisionTreeClassifier import warnings warnings.filterwarnings('ignore')

# Load the dataset df =
pd.read_csv("titanic.csv")

# Drop rows where the target variable is missing
df = df.dropna(subset=['Survived'])

# Select features 'x' and target variable 'y' x =
df[['Pclass', 'Sex', 'Age', 'SibSp', 'Parch', 'Fare']] y =
df["Survived"]

# Encode categorical feature 'Sex' to numeric
le = LabelEncoder() x['Sex'] =
le.fit_transform(x['Sex'])

# Fill missing values in 'Age' with the mean
x['Age'] = x['Age'].fillna(x['Age'].mean())

# Split the dataset into training and testing sets x_train, x_test, y_train, y_test =
train_test_split(x, y, test_size=0.2, random_state=42)

# Create a Random Forest Classifier with 100 decision trees rf_model
= RandomForestClassifier(n_estimators=100, random_state=42)

# Train the Random Forest Classifier
rf_model.fit(x_train, y_train)

# Make predictions using the Random Forest Classifier
y_pred_rf = rf_model.predict(x_test)

# Evaluate the Random Forest Classifier rf_accuracy =
accuracy_score(y_test, y_pred_rf) rf_classification_report =
classification_report(y_test, y_pred_rf)

print("Accuracy of Random Forest Classifier: ", rf_accuracy)
print("Classification Report:\n", rf_classification_report)

# Comparison with other Models
# Initialize models model1 =

KNeighborsClassifier(n_neighbors=9) model2 =
GaussianNB() model3 =
DecisionTreeClassifier(criterion='entropy') model4 =
RandomForestClassifier(n_estimators=100)
# List of models for comparison modellist =
[model1, model2, model3, model4]

# Evaluate each model print("\n=== Model
Comparison Results ===") for model in
modellist:     model.fit(x_train, y_train)
y_pred = model.predict(x_test)
    # Calculate performance metrics
model_accuracy = accuracy_score(y_test, y_pred)
model_confusion_matrix = confusion_matrix(y_test, y_pred)
model_classification_report = classification_report(y_test, y_pred)
    # Display results for each model
print(f"\nModel: {model.__class__.__name__}")
print("Confusion Matrix:")
    print(model_confusion_matrix)
print(f"Accuracy: {model_accuracy:.2f}")
print("Classification Report:")
    print(model_classification_report)

```



Classification Report:		support		precision			
recall	f1-score						
1	0	0.71	0.81	0.76	74	105	
	0.67	0.54	0.60				

accuracy 0.70 179 macro avg 0.69 0.68 0.68 179 weighted
avg 0.69 0.70 0.69 179

Model: GaussianNB

Confusion Matrix:

[[85 20]

[21 53]]

Accuracy: 0.77

Classification Report:

	recall	f1-score	support		precision		
		0	0.80	0.81	0.81	74	105
1		0.73	0.72	0.72			
accuracy					0.77		179
macro avg		0.76	0.76	0.76			179
weighted avg		0.77	0.77	0.77			179

Model: DecisionTreeClassifier

Confusion Matrix:

[[83 22]

[21 53]]

Accuracy: 0.76

Classification Report:

	recall	f1-score	support		precision		
		0	0.80	0.79	0.79	74	105
1		0.71	0.72	0.71			
accuracy					0.76		179
macro avg		0.75	0.75	0.75			179
weighted avg		0.76	0.76	0.76			179

Model: RandomForestClassifier

Confusion Matrix:

[[91 14]

[20 54]]

Accuracy: 0.81

Classification Report:

	recall	f1-score	support		precision		
		0	0.82	0.87	0.84	74	105
1		0.79	0.73	0.76			
accuracy					0.81		179
macro avg		0.81	0.80	0.80			179
weighted avg		0.81	0.81	0.81			179